

# 21377062-王悦扬-第10周作业

---

## 1. Kafaka安装配置

### 1.1 Kafka安装

访问Kafka官网下载页面（<https://kafka.apache.org/downloads>），下载 `kafka_2.13-3.7.0.tgz` 到虚拟机“~/Downloads”目录下。

执行如下命令完成Kafka的安装：

```
1 cd ~/Downloads
2 sudo tar -zxvf kafka_2.13-3.7.0.tgz -C /usr/local
3 cd /usr/local
4 sudo mv kafka_2.13-3.7.0 kafka
5 sudo chown -R hadoop ./kafka
```

```
hadoop@21377062-VirtualBox:~/Downloads$ cd /usr/local
hadoop@21377062-VirtualBox:/usr/local$ sudo mv kafka_2.13-3.7.0 kafka
hadoop@21377062-VirtualBox:/usr/local$ sudo chown -R hadoop ./kafka
```

### 1.2 部署Kafka伪分布式集群

进入kafka目录，在此目录下建立一个目录etc，将config文件夹中的zookeeper.properties复制到etc文件目录中。

```
1 cd /usr/local/kafka
2 sudo mkdir etc
3 cd /usr/local/kafka/config
4 sudo mv zookeeper.properties /usr/local/kafka/etc
```

```

hadoop@21377062-VirtualBox:/usr/local$ cd /usr/local/kafka
hadoop@21377062-VirtualBox:/usr/local/kafka$ sudo mkdir etc
hadoop@21377062-VirtualBox:/usr/local/kafka$ cd /usr/local/kafka/config
hadoop@21377062-VirtualBox:/usr/local/kafka/config$ ls
connect-console-sink.properties    connect-mirror-maker.properties  server.properties
connect-console-source.properties  connect-standalone.properties    tools-log4j.properties
c VBox_GAs_7.0.14.tted.properties  consumer.properties              trogdor.conf
connect-file-sink.properties        kraft                             zookeeper.properties
connect-file-source.properties      log4j.properties
connect-log4j.properties            producer.properties
hadoop@21377062-VirtualBox:/usr/local/kafka/config$ sudo mv zookeeper.properties /usr/local/kafka/etc
hadoop@21377062-VirtualBox:/usr/local/kafka/config$ cd /usr/local/kafka/etc
hadoop@21377062-VirtualBox:/usr/local/kafka/etc$ ls
zookeeper.properties

```

将config文件夹中的server.properties复制三份至etc文件目录中，分别命名为server-0.properties、server-1.properties、server-2.properties

```

1 cd /usr/local/kafka/config
2 sudo cp server.properties /usr/local/kafka/etc/server-0.properties
3 sudo cp server.properties /usr/local/kafka/etc/server-1.properties
4 sudo cp server.properties /usr/local/kafka/etc/server-2.properties

```

```

hadoop@21377062-VirtualBox:/usr/local/kafka/etc$ cd /usr/local/kafka/config
hadoop@21377062-VirtualBox:/usr/local/kafka/config$ sudo cp server.properties /usr/local/kafka/etc/server-0.pr
operties
hadoop@21377062-VirtualBox:/usr/local/kafka/config$ sudo cp server.properties /usr/local/kafka/etc/server-1.pr
operties
hadoop@21377062-VirtualBox:/usr/local/kafka/config$ sudo cp server.properties /usr/local/kafka/etc/server-2.pr
operties
hadoop@21377062-VirtualBox:/usr/local/kafka/config$ cd /usr/local/kafka/etc
hadoop@21377062-VirtualBox:/usr/local/kafka/etc$ ls
server-0.properties  server-1.properties  server-2.properties  zookeeper.properties

```

### 1.3 配置server-X.properties文件

分别编辑三个broker配置server-X.properties文件中的以下信息：

```

1 broker.id = x
2 listeners = PLAINTEXT://:9092(9093,9094)
3 log.dirs.=/tmp/kafka-logsx

```

操作命令：

```

1 cd /usr/local/kafka/etc
2 sudo vim ../etc/server-0.properties
3 sudo vim ../etc/server-1.properties
4 sudo vim ../etc/server-2.properties

```

```

hadoop@21377062-VirtualBox:/usr/local/kafka/config$ cd /usr/local/kafka/etc
hadoop@21377062-VirtualBox:/usr/local/kafka/etc$ ls
s VBox_GAs_7.0.14 ties server-1.properties server-2.properties zookeeper.properties
hadoop@21377062-VirtualBox:/usr/local/kafka/etc$ cd /usr/local/kafka/etc
hadoop@21377062-VirtualBox:/usr/local/kafka/etc$ ls
server-0.properties server-1.properties server-2.properties zookeeper.properties
hadoop@21377062-VirtualBox:/usr/local/kafka/etc$ sudo vim ../etc/server-0.properties
hadoop@21377062-VirtualBox:/usr/local/kafka/etc$ sudo vim ../etc/server-1.properties
hadoop@21377062-VirtualBox:/usr/local/kafka/etc$ sudo vim ../etc/server-2.properties

```

对于3个伪分布式集群分别进行修改

## 1.4 启动zookeeper服务器和kafka集群

首先启动zookeeper:

```

1 cd /usr/local/kafka
2 ./bin/zookeeper-server-start.sh etc/zookeeper.properties

```

```

[2024-05-06 14:26:06,769] INFO Snapshotting: 0x0 to /tmp/zookeeper/version-2/snapshot.0 (org.apache.zookeeper.
server.persistence.FileTxnSnapLog)
[2024-05-06 14:26:06,775] INFO Snapshot loaded in 19 ms, highest zxid is 0x0, digest is 1371985504 (org.apache
.zookeeper.server.ZKDatabase)
[2024-05-06 14:26:06,803] INFO Snapshotting: 0x0 to /tmp/zookeeper/version-2/snapshot.0 (org.apache.zookeeper.
server.persistence.FileTxnSnapLog)
[2024-05-06 14:26:06,804] INFO Snapshot taken in 1 ms (org.apache.zookeeper.server.ZooKeeperServer)
[2024-05-06 14:26:06,823] INFO zookeeper.request_throttler.shutdownTimeout = 10000 ms (org.apache.zookeeper.se
rver.RequestThrottler)
[2024-05-06 14:26:06,823] INFO PrepRequestProcessor (sid:0) started, reconfigEnabled=false (org.apache.zookeep
er.server.PrepRequestProcessor)
[2024-05-06 14:26:06,855] INFO Using checkIntervalMs=60000 maxPerMinute=10000 maxNeverUsedIntervalMs=0 (org.ap
ache.zookeeper.server.ContainerManager)
[2024-05-06 14:26:06,857] INFO ZooKeeper audit is disabled. (org.apache.zookeeper.audit.ZKAuditProvider)

```

启动kafka集群

```

1 cd /usr/local/kafka
2 ./bin/kafka-server-start.sh etc/server-0.properties
3 ./bin/kafka-server-start.sh etc/server-1.properties
4 ./bin/kafka-server-start.sh etc/server-2.properties

```

```
hadoop@21377062-VirtualBox: /usr/local/kafka
[Thunderbird Mail 28:02,445] INFO [ThrottledChannelReaper-Produce]: Shutdown completed (kafka.server.ClientQuotaManager$ThrottledChannelReaper)
[2024-05-06 14:28:02,446] INFO [ThrottledChannelReaper-Request]: Shutting down (kafka.server.ClientQuotaManager$ThrottledChannelReaper)
[2024-05-06 14:28:02,446] INFO [ThrottledChannelReaper-Request]: Stopped (kafka.server.ClientQuotaManager$ThrottledChannelReaper)
[2024-05-06 14:28:02,446] INFO [ThrottledChannelReaper-Request]: Shutdown completed (kafka.server.ClientQuotaManager$ThrottledChannelReaper)
[2024-05-06 14:28:02,446] INFO [ThrottledChannelReaper-ControllerMutation]: Shutting down (kafka.server.ClientQuotaManager$ThrottledChannelReaper)
[2024-05-06 14:28:02,446] INFO [ThrottledChannelReaper-ControllerMutation]: Stopped (kafka.server.ClientQuotaManager$ThrottledChannelReaper)
[2024-05-06 14:28:02,446] INFO [ThrottledChannelReaper-ControllerMutation]: Shutdown completed (kafka.server.ClientQuotaManager$ThrottledChannelReaper)
[2024-05-06 14:28:02,447] INFO [SocketServer listenerType=ZK_BROKER, nodeId=1] Shutting down socket server (kafka.network.SocketServer)
[2024-05-06 14:28:02,495] INFO [SocketServer listenerType=ZK_BROKER, nodeId=1] Shutdown completed (kafka.network.SocketServer)
[2024-05-06 14:28:02,496] INFO Metrics scheduler closed (org.apache.kafka.common.metrics.Metrics)
[2024-05-06 14:28:02,496] INFO Metrics reporters closed (org.apache.kafka.common.metrics.Metrics)
```

启动集群看是否成功

```
1 jps
```

```
hadoop@21377062-VirtualBox: /usr/local/kafka$ jps
3521 Kafka
3042 Kafka
2615 QuorumPeerMain
5336 Kafka
5805 Jps
```

## 2. Java API实现数据统计到JSON

### 2.1 配置Topic

创建3个Topic以完成任务，分别是 `comments`, `likes`, 和 `shares`

```

1 cd /usr/local/kafka
2 ./bin/kafka-topics.sh --create --topic comments --partitions 3 --
  replication-factor 2 --bootstrap-server
  localhost:9092,localhost:9093,localhost:9094
3
4 ./bin/kafka-topics.sh --create --topic likes --partitions 3 --
  replication-factor 2 --bootstrap-server
  localhost:9092,localhost:9093,localhost:9094
5
6 ./bin/kafka-topics.sh --create --topic shares --partitions 3 --
  replication-factor 2 --bootstrap-server
  localhost:9092,localhost:9093,localhost:9094

```

```

hadoop@21377062-VirtualBox:/usr/local/kafka$ cd /usr/local/kafka
hadoop@21377062-VirtualBox:/usr/local/kafka$ ./bin/kafka-topics.sh --create --topic comments --partitions 3 --
replication-factor 2 --bootstrap-server localhost:9092,localhost:9093,localhost:9094
Created topic comments.
hadoop@21377062-VirtualBox:/usr/local/kafka$ ./bin/kafka-topics.sh --create --topic likes --partitions 3 --rep
lication-factor 2 --bootstrap-server localhost:9092,localhost:9093,localhost:9094
Created topic likes.
hadoop@21377062-VirtualBox:/usr/local/kafka$ ./bin/kafka-topics.sh --create --topic shares --partitions 3 --re
plication-factor 2 --bootstrap-server localhost:9092,localhost:9093,localhost:9094
Created topic shares.

```

查看创建的Topics

```

1 cd /usr/local/kafka
2 ./bin/kafka-topics.sh --describe --topic comments --bootstrap-server
  localhost:9092,localhost:9093,localhost:9094
3 ./bin/kafka-topics.sh --describe --topic likes,shares --bootstrap-
  server localhost:9092,localhost:9093,localhost:9094

```

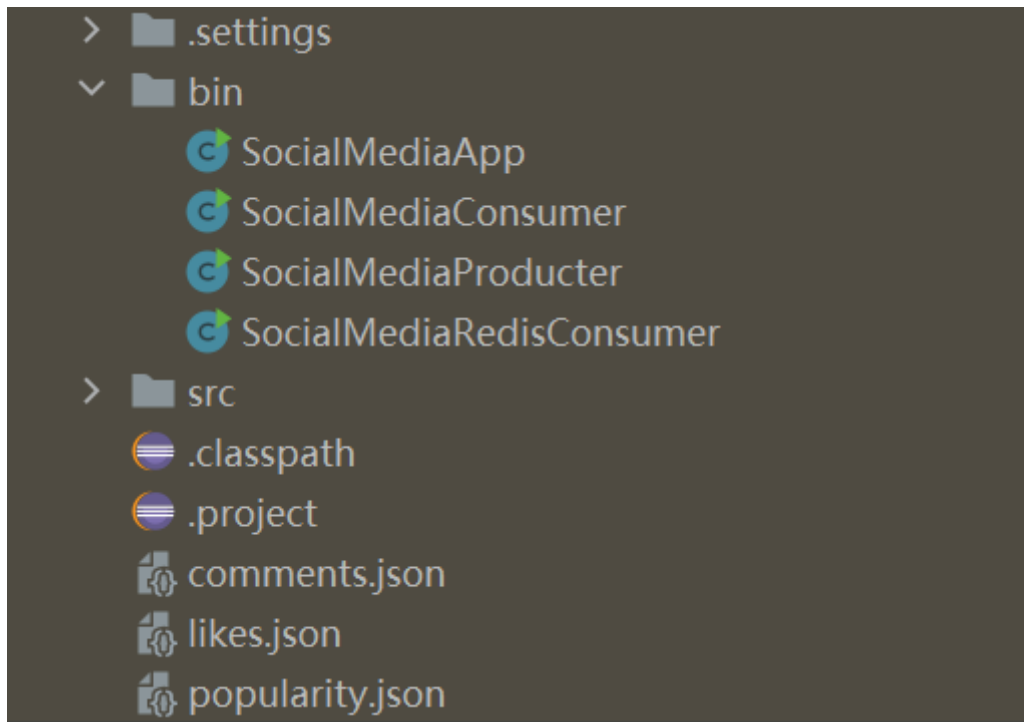
```

hadoop@21377062-VirtualBox:/usr/local/kafka$ cd /usr/local/kafka
hadoop@21377062-VirtualBox:/usr/local/kafka$ ./bin/kafka-topics.sh --describe --topic comments --bootstrap-ser
ver localhost:9092,localhost:9093,localhost:9094
Topic: comments TopicId: 0X_GRQGJT26AzivLznVgNQ PartitionCount: 3 ReplicationFactor: 2 Configs:
Topic: comments Partition: 0 Leader: 1 Replicas: 1,2 Isr: 1,2
Topic: comments Partition: 1 Leader: 0 Replicas: 0,1 Isr: 0,1
Topic: comments Partition: 2 Leader: 2 Replicas: 2,0 Isr: 2,0
h Terminal 377062-VirtualBox:/usr/local/kafka$ ./bin/kafka-topics.sh --describe --topic likes,shares --bootstrap
-server localhost:9092,localhost:9093,localhost:9094
Topic: shares TopicId: onUDO_79QC0Eib81r2LZCA PartitionCount: 3 ReplicationFactor: 2 Configs:
Topic: shares Partition: 0 Leader: 0 Replicas: 0,2 Isr: 0,2
Topic: shares Partition: 1 Leader: 2 Replicas: 2,1 Isr: 2,1
Topic: shares Partition: 2 Leader: 1 Replicas: 1,0 Isr: 1,0
Topic: likes TopicId: Ar2LUKXiR6KA0jdyZmBG4w PartitionCount: 3 ReplicationFactor: 2 Configs:
Topic: likes Partition: 0 Leader: 0 Replicas: 0,1 Isr: 0,1
Topic: likes Partition: 1 Leader: 2 Replicas: 2,0 Isr: 2,0
Topic: likes Partition: 2 Leader: 1 Replicas: 1,2 Isr: 1,2

```

## 2.2 项目创建与环境配置

创建java项目，导入kafka/libs下的所有jar包，并创建3个文件（SocialMediaApp.class, SocialMediaConsumer.class, SocialMediaProducer.class）



## 2.3 程序编写

SocialMediaConsumer.class

```
1 import com.fasterxml.jackson.databind.ObjectMapper;
2 import com.fasterxml.jackson.databind.node.ObjectNode;
3 import java.io.FileWriter;
4 import java.util.Arrays;
5 import java.util.HashMap;
6 import java.util.HashSet;
7 import java.util.Iterator;
8 import java.util.Map;
9 import java.util.Properties;
10 import java.util.Set;
11 import org.apache.kafka.clients.consumer.ConsumerRecord;
12 import org.apache.kafka.clients.consumer.ConsumerRecords;
13 import org.apache.kafka.clients.consumer.KafkaConsumer;
14 import org.apache.kafka.common.serialization.StringDeserializer;
15
16 public class SocialMediaConsumer {
```

```

17     private static final ObjectMapper mapper = new ObjectMapper();
18
19     public SocialMediaConsumer() {
20     }
21
22     public static void main(String[] args) {
23         Properties props = new Properties();
24         props.put("bootstrap.servers",
25 "localhost:9092,localhost:9093,localhost:9094");
26         props.put("group.id", "social-media-group");
27         props.put("key.deserializer",
28 StringDeserializer.class.getName());
29         props.put("value.deserializer",
30 StringDeserializer.class.getName());
31         props.put("auto.offset.reset", "earliest");
32         KafkaConsumer<String, String> consumer = new
33 KafkaConsumer(props);
34         consumer.subscribe(Arrays.asList("likes", "comments",
35 "shares"));
36         Map<String, Set<String>> userComments = new HashMap();
37         Map<String, Map<String, Integer>> userLikes = new
38 HashMap();
39         Map<String, Integer> userPopularity = new HashMap();
40         int giveUp = true;
41         int noRecordsCount = 0;
42
43         try {
44             while(true) {
45                 ConsumerRecords<String, String> records =
46 consumer.poll(100L);
47                 if (records.count() == 0) {
48                     noRecordsCount += 100;
49                     if (noRecordsCount > 10000) {
50                         return;
51                     }
52                 } else {
53                     noRecordsCount = 0;
54                     System.out.println("Received " +
55 records.count() + " records");
56                 }
57
58                 Iterator var10 = records.iterator();
59             }
60         }
61     }
62 }

```

```

51
52         while(var10.hasNext()) {
53             ConsumerRecord<String, String> record =
(ConsumerRecord)var10.next();
54             processRecord(record, userComments, userLikes,
userPopularity);
55         }
56
57         if (!records.isEmpty()) {
58             try {
59                 writeToJsonFiles(userComments, userLikes,
userPopularity);
60             } catch (Exception var14) {
61             }
62         }
63     }
64 } finally {
65     consumer.close();
66     System.out.println("Consumer closed");
67 }
68 }
69
70 private static void processRecord(ConsumerRecord<String,
String> record, Map<String, Set<String>> userComments, Map<String,
Map<String, Integer>> userLikes, Map<String, Integer>
userPopularity) {
71     System.out.println((String)record.value());
72     String[] parts = ((String)record.value()).split(" ");
73     String userWhoPosted = parts[1];
74     String postId = parts[2];
75     System.out.println(topic);
76     switch (topic) {
77         case "shares":
78             int shareCount = parts.length - 3;
79             userPopularity.merge(userWhoPosted, 20 *
shareCount, Integer::sum);
80             break;
81         case "comments":
82             String comment = parts[3];
83             ((Set)userComments.computeIfAbsent(userWhoPosted,
(k) -> {
84                 return new HashSet();

```



```

85         })).add(comment);
86         userPopularity.merge(userWhoPosted, 5,
Integer::sum);
87         break;
88         case "likes":
89             ((Map)userLikes.computeIfAbsent(userWhoPosted, (k)
-> {
90                 return new HashMap();
91             })).put(postId, (Integer)
((Map)userLikes.getDefault(userWhoPosted, new
HashMap()))).getOrDefault(postId, 0) + 1);
92         userPopularity.merge(userWhoPosted, 1,
Integer::sum);
93     }
94
95 }
96
97 private static void writeToJsonFiles(Map<String, Set<String>>
userComments, Map<String, Map<String, Integer>> userLikes,
Map<String, Integer> userPopularity) throws Exception {
98     Throwable var3 = null;
99     Object var4 = null;
100
101     try {
102         FileWriter commentsWriter = new
FileWriter("comments.json");
103
104         try {
105             FileWriter likesWriter = new
FileWriter("likes.json");
106
107             try {
108                 FileWriter popularityWriter = new
FileWriter("popularity.json");
109
110                 try {
111                     mapper.writeValue(commentsWriter,
userComments);
112                     mapper.writeValue(likesWriter, userLikes);
113                     ObjectNode popularityJson =
mapper.createObjectNode();

```

```
114         userPopularity.forEach((user, popularity) -
> {
115             popularityJson.put(user,
(double)popularity / 1000.0);
116         });
117         mapper.writeValue(popularityWriter,
popularityJson);
118     } finally {
119         if (popularityWriter != null) {
120             popularityWriter.close();
121         }
122     }
123 }
124 } catch (Throwable var26) {
125     if (var3 == null) {
126         var3 = var26;
127     } else if (var3 != var26) {
128         var3.addSuppressed(var26);
129     }
130
131     if (likesWriter != null) {
132         likesWriter.close();
133     }
134
135     throw var3;
136 }
137
138     if (likesWriter != null) {
139         likesWriter.close();
140     }
141 } catch (Throwable var27) {
142     if (var3 == null) {
143         var3 = var27;
144     } else if (var3 != var27) {
145         var3.addSuppressed(var27);
146     }
147
148     if (commentsWriter != null) {
149         commentsWriter.close();
150     }
151
152     throw var3;
```

```

153         }
154
155         if (commentsWriter != null) {
156             commentsWriter.close();
157         }
158
159     } catch (Throwable var28) {
160         if (var3 == null) {
161             var3 = var28;
162         } else if (var3 != var28) {
163             var3.addSuppressed(var28);
164         }
165
166         throw var3;
167     }
168 }
169 }

```

#### SocialMediaProducer.class

```

1  import java.io.BufferedReader;
2  import java.io.FileReader;
3  import java.io.IOException;
4  import java.util.Properties;
5  import org.apache.kafka.clients.producer.Callback;
6  import org.apache.kafka.clients.producer.KafkaProducer;
7  import org.apache.kafka.clients.producer.Producer;
8  import org.apache.kafka.clients.producer.ProducerRecord;
9  import org.apache.kafka.clients.producer.RecordMetadata;
10
11  public class SocialMediaProducer {
12      public SocialMediaProducer() {
13      }
14
15      public static void main(String[] args) {
16          String inputFile =
17              "/home/hadoop/Downloads/datasets/student_dataset.txt";
18          Properties props = new Properties();
19          props.put("bootstrap.servers",
20              "localhost:9092,localhost:9093,localhost:9094");

```

[illegible]

```

53         break label295;
54     }
55
56     topic = "shares";
57     break;
58     case 950398559:
59         if
60 (type.equals("comment")) {
61             topic = "comments";
62             break;
63         }
64         default:
65             break label295;
66     }
67     producer.send(new
68 ProducerRecord(topic, (Object)null, message), new Callback() {
69         public void
70 onCompletion(RecordMetadata metadata, Exception exception) {
71             if (exception != null)
72 {
73         System.err.println("Failed to send message: " + message + " to
74 topic: " + topic);
75
76         exception.printStackTrace();
77     } else {
78         System.out.println("Sent message: " + message + " to topic: " +
79 topic);
80     }
81 }
82     });
83     continue;
84 }
85
86 System.err.println("Unknown type: "
87 + type + " in line: " + line);
88 }
89 }

```

```

85         return;
86     }
87     } finally {
88         if (br != null) {
89             br.close();
90         }
91     }
92 }
93 } catch (Throwable var27) {
94     if (var4 == null) {
95         var4 = var27;
96     } else if (var4 != var27) {
97         var4.addSuppressed(var27);
98     }
99
100     throw var4;
101 }
102 } catch (IOException var28) {
103     var28.printStackTrace();
104 } finally {
105     producer.close();
106 }
107
108 }
109 }

```

## SocialMediaApp.class

```

1  import java.io.IOException;
2
3  public class SocialMediaApp {
4      public static void main(String[] args) {
5          Thread consumerThread = new Thread(() -> {
6              SocialMediaRedisConsumer.main(new String[]{});
7          });
8
9          Thread producerThread = new Thread(() -> {
10             try {
11                 Thread.sleep(500);
12             } catch (InterruptedException e) {
13                 e.printStackTrace();

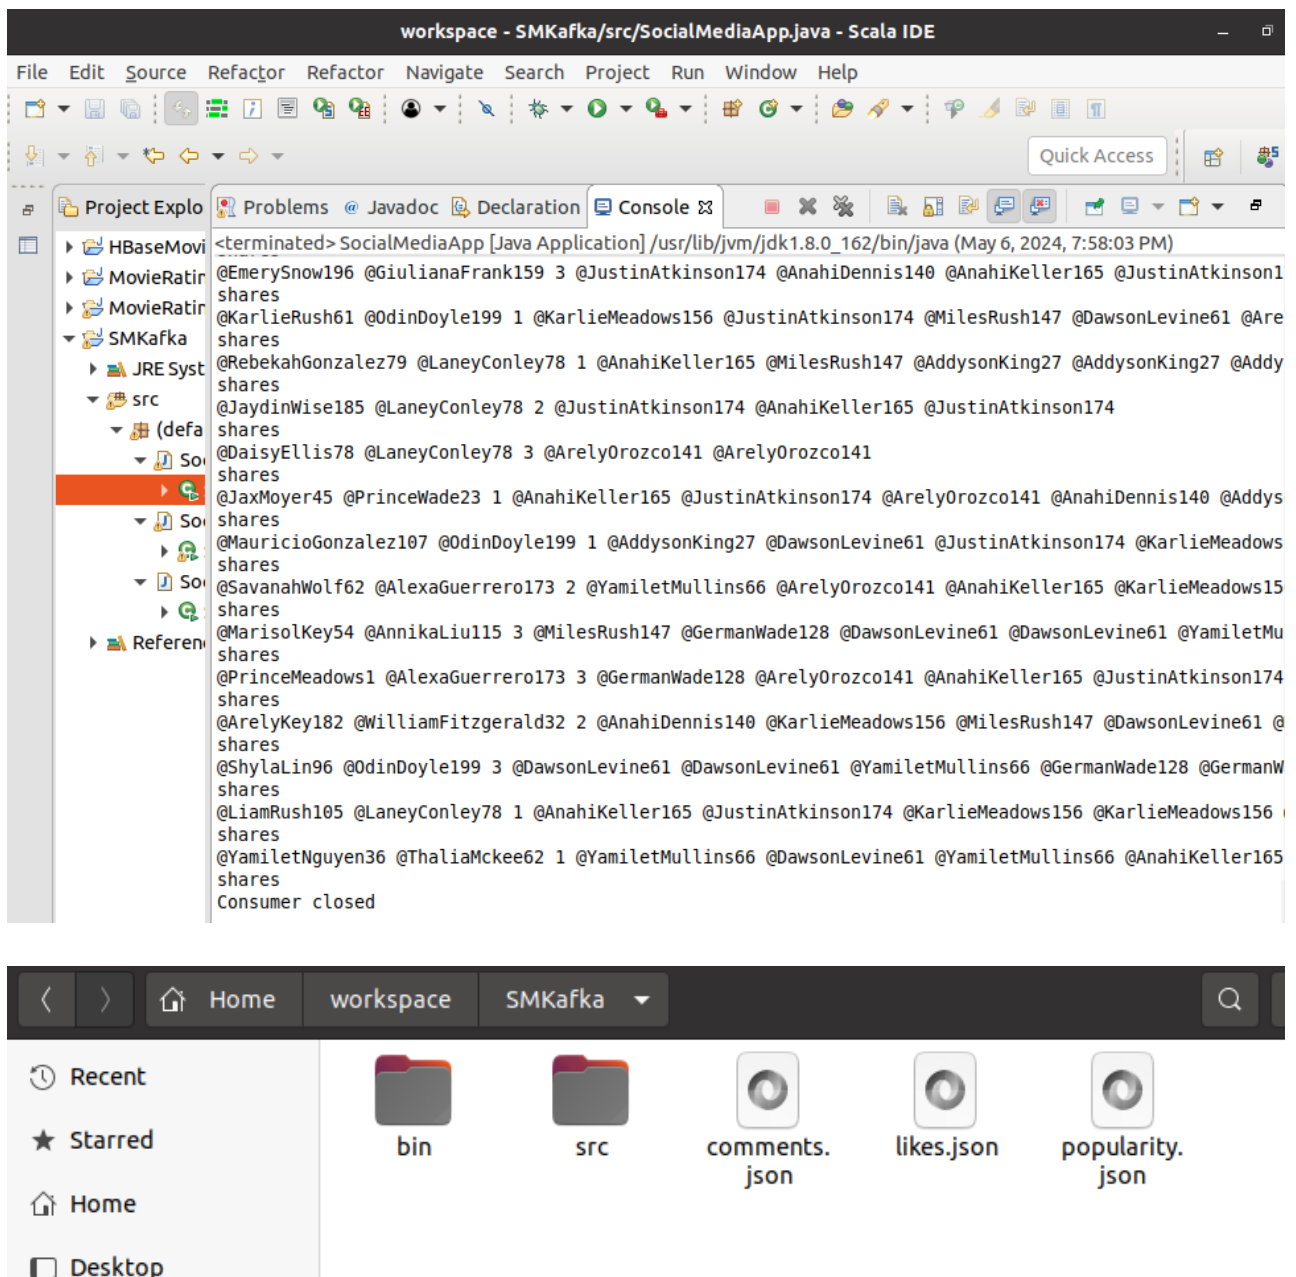
```

```

14         }
15         SocialMediaProducer.main(new String[]{});
16     });
17
18     consumerThread.start();
19     producerThread.start();
20 }
21 }

```

## 2.4 运行结果



## 3. Java API数据统计到Redis

### 3.1 安装配置Redis数据库

```
1 sudo apt-get install redis-server
```

```
Setting up libjemalloc2:amd64 (5.2.1-1ubuntu1) ...
Setting up lua-cjson:amd64 (2.1.0+dfsg-2.1) ...
Setting up lua-bitop:amd64 (1.0.2-5) ...
Setting up liblua5.1-0:amd64 (5.1.5-8.1build4) ...
Setting up libhiredis0.14:amd64 (0.14.0-6) ...
Setting up redis-tools (5:5.0.7-2ubuntu0.1) ...
Setting up redis-server (5:5.0.7-2ubuntu0.1) ...
Created symlink /etc/systemd/system/redis.service → /lib/systemd/system/redis-server.service.
Created symlink /etc/systemd/system/multi-user.target.wants/redis-server.service → /lib/systemd/system/redis-server.service.
Processing triggers for systemd (245.4-4ubuntu3.23) ...
Processing triggers for man-db (2.9.1-1) ...
Processing triggers for libc-bin (2.31-0ubuntu9.14) ...
```

查看redis是否正常启动

```
1 service redis-server status
```

```
hadoop@21377062-VirtualBox:~$ service redis-server status
● redis-server.service - Advanced key-value store
   Loaded: loaded (/lib/systemd/system/redis-server.service; enabled; vendor preset: enabled)
   Active: active (running) since Mon 2024-05-06 19:01:56 CST; 1h 5min ago
     Docs: http://redis.io/documentation,
           man:redis-server(1)
   Main PID: 33265 (redis-server)
     Tasks: 4 (limit: 9429)
    Memory: 1.8M
    CGroup: /system.slice/redis-server.service
           Show Applications 3265 /usr/bin/redis-server 127.0.0.1:6379
```

### 3.2 程序编写

SocialMediaRedisConsumer.class

```
1 import redis.clients.jedis.Jedis;
2 import org.apache.kafka.clients.consumer.ConsumerRecord;
3 import org.apache.kafka.clients.consumer.ConsumerRecords;
4 import org.apache.kafka.clients.consumer.KafkaConsumer;
5 import org.apache.kafka.clients.consumer.ConsumerConfig;
6 import org.apache.kafka.common.serialization.StringDeserializer;
7
8 import java.util.*;
9
10 public class SocialMediaRedisConsumer {
```



```

11     private static Jedis jedis;
12
13     public static void main(String[] args) {
14         jedis = new Jedis("localhost", 6379);
15         System.out.println("Connected to Redis");
16
17         Properties props = new Properties();
18         props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG,
19 "localhost:9092,localhost:9093,localhost:9094");
20         props.put(ConsumerConfig.GROUP_ID_CONFIG, "social-media-
21 group");
22         props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
23 StringDeserializer.class.getName());
24         props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
25 StringDeserializer.class.getName());
26         props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG,
27 "earliest");
28
29         KafkaConsumer<String, String> consumer = new KafkaConsumer<>
30 (props);
31         consumer.subscribe(Arrays.asList("likes", "comments",
32 "shares"));
33
34         final int giveUp = 10000; // 10 seconds timeout
35         int noRecordsCount = 0;
36
37         try {
38             while (true) {
39                 ConsumerRecords<String, String> records =
40 consumer.poll(100);
41                 if (records.count() == 0) {
42                     noRecordsCount += 100;
43                     if (noRecordsCount > giveUp) break;
44                 } else {
45                     noRecordsCount = 0;
46                     System.out.println("Received " + records.count()
47 + " records");
48                 }
49
50                 for (ConsumerRecord<String, String> record :
51 records) {
52                     processRecord(record);
53                 }
54             }
55         }
56     }
57 }

```

```

43         }
44     }
45     } finally {
46         consumer.close();
47         jedis.close();
48         System.out.println("Consumer and Redis client closed");
49     }
50 }
51
52 private static void processRecord(ConsumerRecord<String, String>
record) {
53     String[] parts = record.value().split(" ");
54     String topic = record.topic();
55     String userWhoPosted = parts[1];
56     String postId = parts[2];
57
58     switch (topic) {
59         case "likes":
60             String likeskey = "likes:" + userWhoPosted + ":" +
postId;
61             jedis.hincrBy(userWhoPosted, postId, 1);
62             break;
63         case "comments":
64             String comment = parts[3];
65             jedis.rpush("comments:" + userWhoPosted, comment);
66             break;
67         case "shares":
68             int shareCount = parts.length - 3;
69             jedis.incrBy("popularity:" + userWhoPosted, 20 *
shareCount);
70             break;
71     }
72 }
73 }

```

### 3.3 运行结果

```
<terminated> SocialMediaApp [Java Application] /usr/lib/jvm/jdk1.8.0_162/bin/java (May 6, 2024, 8:10:49 PM)
Sent message: @MarisolKey54 @AnnikaLiu115 3 "vsLiWl " to topic: comments
Sent message: @PrinceMeadows1 @AlexaGuerrero173 3 "oQJhHUGDUwCZVoFCTLWtnpTq TixwSnUYnllBFBAImmIO UIQvYoKmPqL HbIRQGQeAMSwqWIX
Sent message: @ArellyKey182 @WilliamFitzgerald32 2 "HdKggygfvxnarNTk KDTSJIAIAyhNGKw vNLzLTzamyNYtygyvFfBtbrR " to topic: comme
Sent message: @ShylaLin96 @OdinDoyle199 3 "bbkqfKHhOrp MNFGtMBHwdKHqrASZEIYNFMv iMmvSDcpGgWOL " to topic: comments
Sent message: @LiamRush105 @LaneyConley78 1 "rwcBRyUWxOZqptM NIWMJbzX Okv0s jvdvVcLMrLIAwXZ " to topic: comments
Sent message: @YamiletNguyen36 @ThaliaMckee62 1 "ldGBEEBMreowPlet nIgxLtFbLYCaxgf tyjMMKXgaIppqmfRjRxHhunR JDLeFmrQGUXg " to t
Sent message: @DelaneyMorris61 @LaneyConley78 2 to topic: likes
Sent message: @VeronicaBruce37 @AlexaGuerrero173 3 to topic: likes
Sent message: @EnriqueMorrow21 @PrinceKing174 3 to topic: likes
Sent message: @LiamStephenson62 @AlexaGreen177 1 to topic: likes
Sent message: @AzulDeleon137 @OdinDoyle199 1 to topic: likes
Sent message: @GradyMckee131 @PrinceKing174 1 to topic: likes
Sent message: @JaydinRussell77 @WilliamFitzgerald32 1 to topic: likes
Sent message: @RhysBrooks25 @PrinceWade23 2 to topic: likes
Sent message: @MosheHuynh12 @OdinDoyle199 3 to topic: likes
Sent message: @GiselleMeadows58 @AlexaGreen177 2 to topic: likes
Sent message: @MosheAtkinson53 @AlexaGuerrero173 3 to topic: likes
Sent message: @DominickLevine117 @WilliamFitzgerald32 1 to topic: likes
Sent message: @KaleyStevenson112 @OdinDoyle199 2 to topic: likes
Sent message: @DannaKey152 @AlexaGreen177 3 to topic: likes
Sent message: @EmerySnow196 @GiulianaFrank159 3 to topic: likes
Sent message: @KarlieRush61 @OdinDoyle199 1 to topic: likes
Sent message: @RebekahGonzalez79 @LaneyConley78 1 to topic: likes
Sent message: @JaydinWise185 @LaneyConley78 2 to topic: likes
Sent message: @DaisyEllis78 @LaneyConley78 3 to topic: likes
Sent message: @JaxMoyer45 @PrinceWade23 1 to topic: likes
Sent message: @MauricioGonzalez107 @OdinDoyle199 1 to topic: likes
Sent message: @SavanahWolf62 @AlexaGuerrero173 2 to topic: likes
Sent message: @MarisolKey54 @AnnikaLiu115 3 to topic: likes
Sent message: @PrinceMeadows1 @AlexaGuerrero173 3 to topic: likes
Sent message: @ArellyKey182 @WilliamFitzgerald32 2 to topic: likes
Sent message: @ShylaLin96 @OdinDoyle199 3 to topic: likes
Sent message: @LiamRush105 @LaneyConley78 1 to topic: likes
Sent message: @YamiletNguyen36 @ThaliaMckee62 1 to topic: likes
Received 8 records
Received 99 records
Consumer and Redis client closed
```

查看运行结果

```
1 redis-cli
2 KEYS *
```

```

hadoop@21377062-VirtualBox:~/workspace$ cd SocialMediaKa
hadoop@21377062-VirtualBox:~/workspace/SocialMediaKa$ redis-cli
127.0.0.1:6379> KEYS *
 1) "@ThaliaMckee62"
 2) "@LaneyConley78"
 3) "popularity:@PrinceKing174"
 4) "popularity:@LaneyConley78"
 5) "comments:@OdinDoyle199"
 6) "popularity:@OdinDoyle199"
 7) "comments:@AlexaGuerrero173"
 8) "comments:@WilliamFitzgerald32"
 9) "comments:@AlexaGreen177"
10) "@WilliamFitzgerald32"
11) "@PrinceKing174"
12) "popularity:@AlexaGreen177"
13) "@OdinDoyle199"
14) "popularity:@PrinceWade23"
15) "@PrinceWade23"
16) "popularity:@ThaliaMckee62"
17) "popularity:@AlexaGuerrero173"
18) "comments:@ThaliaMckee62"
19) "comments:@AnnikaLiu115"
20) "@AlexaGuerrero173"
21) "@GiulianaFrank159"
22) "comments:@LaneyConley78"
23) "@AlexaGreen177"
24) "popularity:@GiulianaFrank159"
25) "comments:@GiulianaFrank159"
26) "popularity:@AnnikaLiu115"
27) "popularity:@WilliamFitzgerald32"
28) "comments:@PrinceWade23"
29) "comments:@PrinceKing174"
30) "@AnnikaLiu115"

```

## 4. 问题及解决方案

### 4.1 启动Kafka集群时报错

检查发现配置文件时未修改broker id，但修改完成后仍然报错：

```

[2024-05-06 14:31:49,712] ERROR Exiting Kafka due to fatal exception during startup. (kafka.Kafka$)
java.lang.RuntimeException: Stored node id 1 doesn't match previous node id 3 in /tmp/kafka-logs2/meta.properties. If you moved your data, make sure your configured node id matches. If you intend to create a new node, you should remove all data in your data directories.
    at org.apache.kafka.metadata.properties.MetaPropertiesEnsemble.verify(MetaPropertiesEnsemble.java:526)
    at kafka.server.KafkaServer.startup(KafkaServer.scala:250)
    at kafka.Kafka$.main(Kafka.scala:112)
    at kafka.Kafka.main(Kafka.scala)
[2024-05-06 14:31:49,715] INFO shutting down (kafka.server.KafkaServer)

```

查询日志文件的存储位置：

```
# A comma separated list of directories under which to store log files
log.dirs=/tmp/kafka-logs2
```

```
hadoop@21377062-VirtualBox:/usr/local/kafka/bin$ sudo find / -name "kafka-logs2"
[sudo] password for hadoop:
find: '/run/user/1001/gvfs': Permission denied
/tmp/kafka-logs2
hadoop@21377062-VirtualBox:/usr/local/kafka/bin$ cd /tmp/kafka-logs2
hadoop@21377062-VirtualBox:/tmp/kafka-logs2$ ll
total 24
drwxrwxr-x  2 hadoop hadoop 4096 5月  6 14:28 ./
drwxrwxrwt 22 root   root   4096 5月  6 14:28 ../
-rw-rw-r--  1 hadoop hadoop   0 5月  6 14:28 cleaner-offset-checkpoint
-rw-rw-r--  1 hadoop hadoop  30 5月  6 14:28 .kafka_cleanshutdown
-rw-rw-r--  1 hadoop hadoop   0 5月  6 14:28 .lock
-rw-rw-r--  1 hadoop hadoop   4 5月  6 14:28 log-start-offset-checkpoint
-rw-rw-r--  1 hadoop hadoop  88 5月  6 14:28 meta.properties
-rw-rw-r--  1 hadoop hadoop   4 5月  6 14:28 recovery-point-offset-checkpoint
-rw-rw-r--  1 hadoop hadoop   0 5月  6 14:28 replication-offset-checkpoint
hadoop@21377062-VirtualBox:/tmp/kafka-logs2$ vim meta.properties
```

编辑meta.properties的broker id，与配置文件一致后，成功解决报错。